

A Multi-Dimensional Framework for Ethical and Environmental Evaluation of Generative AI Systems

Bugatha Ganasyam Mani Sankar

Department of Computer Science and Engineering, JNTU-GV College of Engineering, Vizianagaram, India
ganasyammanisankarbugatha@gmail.com

Abstract---

This paper proposes a multi-dimensional framework for evaluating the ethical risks and environmental sustainability of Generative AI systems. We introduce a Responsible AI (RAI) scoring model that integrates Ethical Score (E), Environmental Score (S), and Transparency Score (T) using weighted aggregation. Experimental evaluation across simulated tactical scenarios demonstrates that the proposed framework identifies ethical risks—including subtle demographic biases undetectable by surface-level fairness metrics—with substantially greater sensitivity than keyword-based baseline approaches. The framework is designed for practical deployment in high-stakes domains including military simulations, policy recommendations, and crisis response systems. This research contributes to the emerging field of Green AI [19] by providing actionable metrics for sustainable AI governance aligned with UN Sustainable Development Goals and ISO/IEC TR 20226:2025 standards.

Index Terms---Generative AI, Ethical AI, Green AI, Responsible AI, Environmental Impact, Bias Detection, Sustainability, Carbon Footprint, Water Footprint, AI Governance

I. INTRODUCTION

Generative Artificial Intelligence (GenAI) has transitioned from a research curiosity to a transformative force reshaping industries, governance, and daily human interaction. Large Language Models (LLMs) such as GPT-4, Claude, and Gemini, alongside diffusion-based image generators and video synthesis models, now power applications ranging from creative content production to strategic military simulations and public policy formulation [1][21]. The GPT-4 Technical Report [1] documents the model's deployment across professional and academic benchmarks, while recent surveys highlight its expanding use in high-stakes decision-making contexts [21].

The ethical dimension of Generative AI encompasses algorithmic bias, lack of accountability in high-stakes decision-making, and the propagation of harmful content. In military and strategic contexts, biased AI recommendations can lead to catastrophic outcomes, while in policy domains, opaque decision-making undermines democratic accountability [2]. Concurrently, the environmental toll of training and deploying these models has reached alarming proportions. A single training run of GPT-3 emitted approximately 626,000 pounds of CO₂—equivalent to 300 round-trip flights between New York and San Francisco [3]. The AI boom of 2025 alone released roughly as much CO₂ into the atmosphere as New York City, with annual emissions exceeding 50 million metric tonnes [4].

Recent findings from the United Nations University Institute for Water, Environment and Health (UNU-INWEH) reveal that AI-related water consumption could equal the basic annual domestic needs of 1.3 billion people by 2030, while the land footprint of data centers may exceed 14,500 square kilometers—twice the size of the Jakarta metropolitan area [5]. Furthermore, the International Energy Agency (IEA) projects that AI data center energy demand will more than quadruple by 2030, with only approximately half of this demand likely to be met by renewable sources [6].

Despite these converging crises, the AI research community has largely treated ethical evaluation and environmental sustainability as orthogonal concerns. This paper bridges this critical gap by proposing a unified, multi-dimensional framework that simultaneously evaluates ethical risks and environmental costs. Our contributions are threefold: (1) a novel ethical risk modeling methodology combining NLP-based bias detection with decision risk classification; (2) an environmental impact estimation system integrating estimated energy, water, and carbon metrics derived from published hardware specifications and grid intensity data; and (3) a Responsible AI (RAI) scoring model that aggregates these dimensions into an actionable governance metric.

II. RELATED WORK

The intersection of ethics and sustainability in AI has garnered increasing scholarly attention, yet comprehensive frameworks remain scarce. Bender et al. [7] introduced the concept of 'stochastic parrots' to critique the environmental and ethical costs of large language models, emphasizing the disproportionate burden on marginalized communities. Patterson et al. [8] provided detailed carbon footprint estimates for Google's ML workloads, establishing foundational metrics for energy-aware AI development.

In the ethical domain, Mehrabi et al. [9] surveyed bias and fairness in machine learning, categorizing sources of bias from data collection through model deployment. Arrieta et al. [10] proposed explainable AI (XAI) techniques to enhance transparency in high-stakes applications. However, these works predominantly address either ethics or sustainability in isolation.


Recent regulatory developments have catalyzed integrated approaches. The European Union's AI Act (2024) mandates risk-based classification of AI systems with explicit environmental reporting requirements [11]. The ISO/IEC TR 20226:2025 standard provides a comprehensive framework for environmental sustainability aspects of AI systems, including workload, resource utilization, carbon impact, and water footprint metrics [12]. Concurrently, the UN Secretary-General's 2025 initiative calls for AI companies to disclose full environmental impacts and commit to 100% renewable energy for data centers by 2030 [13].

Our work extends these foundations by developing an operational framework that translates regulatory and standardization requirements into measurable, actionable metrics. Unlike prior approaches that focus exclusively on training-phase emissions, our framework accounts for the full lifecycle—including inference, which constitutes 80-90% of total energy demand according to UNU-INWEH findings [5].

III. PROPOSED METHODOLOGY

The proposed framework comprises four integrated modules: Ethical Risk Modeling, Tactical Scenario Simulation, Environmental Impact Estimation, and the Responsible AI Scoring Model. The system workflow follows a sequential pipeline: Scenario Input → AI Output Generation → Ethical Analysis → Resource Estimation → Scoring → Optimization Recommendations.

Proposed Framework: System Workflow Pipeline



[Fig. 1. System workflow pipeline: Scenario Input → AI Output Generation → Ethical Analysis → Resource Estimation → RAI Scoring → Optimization Recommendations. The framework operates as a sequential evaluation pipeline with feedback loops for iterative improvement.]

The operational pipeline, depicted in Fig. 1, transforms raw scenario inputs into actionable governance recommendations through six sequential stages, with each module feeding forward into the composite RAI scoring function.

A. Ethical Risk Modeling

The ethical evaluation module quantifies three primary dimensions: Bias Score (B), Explainability Score (X), and Decision Risk Factor (D). The Bias Score is computed using a multi-metric approach combining demographic parity difference, equalized odds, and disparate impact ratio. For NLP outputs, we employ sentiment analysis divergence across protected attributes (gender, ethnicity, age) using pre-trained fairness-aware embeddings.

The Explainability Score leverages attention-weight analysis and LIME (Local Interpretable Model-agnostic Explanations) to quantify the transparency of AI-generated decisions. Higher attention concentration on relevant features correlates with improved explainability. The

Decision Risk Factor employs a hierarchical classification system categorizing outputs into: (1) Low Risk (informational queries), (2) Medium Risk (recommendations with human oversight), and (3) High Risk (autonomous decisions in military, medical, or financial domains).

Formally, the Ethical Score E is defined as:

$$E = \alpha_1(1 - B_norm) + \alpha_2X + \alpha_3(1 - D_norm)$$

where B_norm and D_norm are normalized bias and risk scores, respectively, and $\alpha_1, \alpha_2, \alpha_3$ are weighting coefficients satisfying $\alpha_1 + \alpha_2 + \alpha_3 = 1$. Based on empirical validation across 150 tactical scenarios (50 per category), we set $\alpha_1 = 0.4$, $\alpha_2 = 0.35$, and $\alpha_3 = 0.25$, reflecting the primacy of bias mitigation in ethical AI governance.

B. Environmental Impact Estimation

The environmental module quantifies three critical resource dimensions: Energy Consumption (kWh), Water Usage (liters), and Carbon Footprint (kgCO₂e). Drawing on the UNU-INWEH multi-footprint methodology [5], we distinguish between direct operational impacts and indirect supply-chain effects.

Energy estimation employs the Power Usage Effectiveness (PUE) metric, where $PUE = \text{Total Facility Energy} / \text{IT Equipment Energy}$. Industry-standard AI data centers exhibit PUE values between 1.2 and 1.8, with hyperscale facilities approaching 1.1. For a given inference task, energy consumption E_inf is modeled as:

$$E_inf = P_gpu \times t_inf \times N_gpu \times PUE$$

where P_gpu is the thermal design power per GPU (typically 300-700W for NVIDIA A100/H100), t_inf is inference time, and N_gpu is the number of active GPUs.

Water footprint estimation follows the approach of Li et al. [14], accounting for both direct cooling water and indirect water embedded in electricity generation. The Water Usage Effectiveness (WUE) metric quantifies liters consumed per kWh. For evaporative cooling systems in water-stressed regions, WUE can exceed 2.0 L/kWh, while air-cooled systems achieve $WUE < 0.5$ L/kWh. Recent data indicates that a single ChatGPT conversation of 20-50 exchanges consumes approximately 0.5 liters of fresh water [14].

Carbon footprint calculation integrates regional grid carbon intensity (kgCO₂e/kWh) with energy consumption. Given that 56% of data center electricity derives from fossil fuels [16], the carbon intensity varies dramatically by geographic location---from approximately 50 gCO₂e/kWh in regions with high renewable penetration to over 800 gCO₂e/kWh in coal-dependent grids. This geographic dependency underscores the environmental justice dimension of AI deployment.

C. Responsible AI Scoring Model

The Responsible AI (RAI) score aggregates ethical and environmental performance into a unified governance metric. The model incorporates three component scores: Ethical Score (E), Environmental Score (S), and Transparency Score (T).

The Environmental Score S is computed as the harmonic mean of normalized energy, water, and carbon efficiency metrics:

$$S = 3 / (1/S_e + 1/S_w + 1/S_c)$$

where S_e , S_w , and S_c represent normalized scores for energy efficiency, water efficiency, and carbon efficiency, respectively. The harmonic mean penalizes poor performance in any single dimension, ensuring balanced sustainability.

The Transparency Score T evaluates data provenance, model documentation, and auditability. It incorporates: (1) training data disclosure (0-40 points), (2) model architecture transparency (0-30 points), and (3) environmental reporting compliance (0-30 points), aligned with ISO/IEC TR 20226:2025 requirements.

The composite RAI score is defined as:

$$RAI = w_1E + w_2S + w_3T$$

The weighting coefficients were set based on a systematic review of responsible AI literature and regulatory priorities. The highest weight ($w_1 = 0.4$) is assigned to the Ethical dimension, consistent with Mehrabi et al.'s [9] finding that bias mitigation is the primary concern in high-stakes AI deployment, and with the EU AI Act's [11] risk-based classification that prioritizes ethical harm prevention. The Environmental dimension ($w_2 = 0.35$) reflects the urgency established in UNU-INWEH [5] and ISO/IEC TR 20226:2025 [12] frameworks, while the Transparency dimension ($w_3 = 0.25$) aligns with Arrieta et al.'s [10] XAI taxonomy and emerging disclosure mandates. Empirical calibration of these weights through structured expert consultation (e.g., Delphi method or AHP) is identified as future work in Section VII.

D. Tactical Scenario Simulation

To validate the framework, we constructed three high-stakes scenario categories: (1) Military Strategic Decision-Making, (2) Crisis Resource Allocation, and (3) Autonomous Policy Recommendation. Each scenario presents an AI system with complex ethical dilemmas and resource-intensive computational requirements.

In the military domain, scenarios involve target prioritization, collateral damage assessment, and rules-of-engagement interpretation. The crisis response scenarios simulate natural disaster resource distribution across demographically diverse populations, testing for allocative fairness. Policy scenarios require AI systems to generate regulatory recommendations with environmental impact assessments.

The framework is designed to evaluate outputs from state-of-the-art LLMs including GPT-4, Claude-3, and Llama-3, applying the ethical and environmental evaluation pipelines to generated responses. This multi-model comparison reveals significant variance in both ethical reliability and resource consumption across architectures, as estimated from published model specifications and inference benchmarks.


IV. EXPERIMENTAL RESULTS AND ANALYSIS

A simulation-based experimental evaluation was conducted across 150 tactical scenarios (50 per category) using three state-of-the-art LLMs (GPT-4, Claude-3, Llama-3). All resource estimates were derived from published hardware specifications---specifically, NVIDIA A100 GPU thermal design power (400W), industry-standard Power Usage Effectiveness (PUE = 1.4), and U.S. national average grid carbon intensity (450 gCO₂e/kWh)---combined with literature-derived empirical parameters for inference latency and batch processing [5][8][16]. While absolute figures should be interpreted as model-estimated values rather than live hardware measurements, the relative comparisons across models, tasks, and scenarios remain valid for framework validation purposes.

A. Ethical Risk Detection Performance

To establish a comparative baseline, we implemented a keyword-based bias detection filter that flags outputs containing predefined demographic terms (e.g., race, gender, age descriptors) and applies a simple threshold classifier for risk categorization. This baseline represents the minimal viable ethical oversight currently deployed in many production AI systems [9]. Against this baseline, the proposed multi-metric framework---incorporating demographic parity, equalized odds, disparate impact, and semantic harm detection---achieves substantially improved sensitivity to subtle biases, with the mean Ethical Score (E) rising from 53.7 (baseline) to 66.6 (proposed) on the simulated scenario set. It should be noted that these figures are derived from our simulation and require empirical validation on live systems.

As shown in Fig. 2, military scenarios consistently exhibit the highest risk factors (D = 0.85) and lowest ethical scores (E = 59.8), reflecting the inherent complexity of autonomous decision-making in conflict contexts. Crisis response scenarios achieve more favorable ethical balances (E = 70.8) due to lower baseline risk, while policy recommendations---despite moderate bias---score highest overall (E = 69.3) owing to strong explainability.



[Fig. 2. Ethical evaluation metrics across scenario categories: (a) Ethical Score (E), (b) Bias Score (B), (c) Explainability Score (X), and (d) Decision Risk Factor (D). Military scenarios exhibit the highest risk and lowest ethical scores, while policy scenarios demonstrate the most favorable balance.]

TABLE I


ETHICAL EVALUATION RESULTS ACROSS SCENARIO CATEGORIES

Scenario Type	Bias Score (B)	Explainability (X)	Risk Factor (D)	Ethical Score (E)
Military	0.22	0.71	0.85	59.8
Crisis Response	0.15	0.78	0.62	70.8
Policy	0.18	0.65	0.45	69.3
Average	0.18	0.71	0.64	66.6

B. Environmental Impact Quantification

Table II details the environmental footprint per 1,000 inferences across models and modalities, calculated using a Water Usage Effectiveness (WUE) of 2.0 L/kWh representative of evaporative cooling systems in water-stressed regions. Text generation exhibited the lowest impact (0.14 kWh, 0.28 L water, 63 gCO₂e per 1,000 queries), while image generation consumed approximately 50× more energy (7.2 kWh, 14.4 L water, 3,240 gCO₂e) compared to Claude-3 text inference (0.14 kWh per 1,000 queries). Video generation, estimated for Sora-like architectures, demonstrated the highest per-unit impact at 1.0 kWh per single video, 4.0 L water, and 450 gCO₂e---consistent with recent industry analyses [17]. Note that video figures represent per-video estimates, not per-1,000-inference totals as used for text and image modalities.

Fig. 3 illustrates the dramatic variance in environmental cost across modalities, with image generation requiring 7.2 kWh per 1,000 inferences compared to 0.14 kWh for text---a disparity that underscores the need for modality-aware sustainability metrics. Video generation, estimated at 1.0 kWh per single video, represents the highest per-unit impact among evaluated tasks.



[Fig. 3. Environmental footprint per 1,000 inferences across models (GPT-4, Claude-3, Llama-3) and tasks (text, image, video): (a) Energy consumption (kWh), (b) Water usage (liters), (c) Carbon footprint (gCO₂e). Image generation consumes approximately 50× more energy than text-based inference (Note: values shown per 1,000 inferences.)]

TABLE II

ENVIRONMENTAL FOOTPRINT PER 1,000 INFERENCE


Model / Task	Energy (kWh)	Water (L)	Carbon (gCO ₂ e)	RAI Score
GPT-4 (Text)	0.18	0.36	81	78.2
Claude-3 (Text)	0.14	0.28	63	82.1
Llama-3 (Text)	0.09	0.18	41	73.9
Image Generation	7.20	14.40	3,240	42.6

The framework's optimization recommendations---model quantization (INT8), dynamic batching, and geographic load shifting to renewable grids---are projected to achieve substantial reductions in carbon footprint and water consumption based on published efficiency studies. Quantization alone is estimated to reduce energy per inference by approximately 40% with negligible accuracy degradation (<1% on ethical classification tasks) [8][14][15].

C. RAI Score Validation

The composite RAI scores exhibit internal consistency with their constituent dimensions: models scoring in the Excellent tier (RAI ≥ 85) consistently demonstrated low bias (B < 0.15), efficient resource utilization (PUE < 1.3), and full transparency documentation. Models in the Good tier (RAI 70--84), such as Llama-3 in our simulated evaluation, showed acceptable environmental efficiency but moderate ethical scores due to higher residual bias. Conversely, models scoring < 40 (Unacceptable) showed either severe ethical deficiencies (B > 0.4) or extreme environmental costs (>10 kWh per 1,000 inferences). The scoring thresholds (Excellent ≥ 85, Good 70--84, Moderate 55--69, Poor 40--54, Unacceptable < 40) were aligned with EU AI Act [11] risk tiers and ISO/IEC TR 20226:2025 [12] compliance levels. External validation of the RAI score against domain-expert judgments is identified as a critical direction for future work (Section VII).


The RAI score distribution, visualized in Fig. 4, reveals that no single model dominates across all scenario types. Claude-3 achieves the highest scores in policy and crisis contexts (RAI = 76.4--82.1) due to superior environmental efficiency,, while all models struggle in military contexts (RAI < 65), confirming the necessity of multi-objective optimization in responsible AI deployment.



[Fig. 4. Responsible AI (RAI) score heatmap across scenario types (military, crisis, policy) and models (GPT-4, Claude-3, Llama-3). Scores range from 0 (Unacceptable) to 100 (Excellent). Llama-3 achieves the highest scores in low-risk scenarios, while all models struggle in military contexts.]

Fig. 5(a) shows that only 15% of simulated evaluations achieved Excellent tier (RAI ≥ 85), while the majority (65%) fell into Moderate or Good categories---indicating substantial room for improvement across the industry. Fig. 5(b) projects substantial growth in annual AI

environmental impact by 2030 if current trends continue, with carbon, energy, and water footprints all increasing proportionally with demand [5][6].




[Fig. 5. (a) RAI score classification distribution across 150 simulated scenarios; (b) Projected annual AI environmental impact for 2025 and 2030 at scale (1M daily queries), showing energy (TWh), water (billion liters), and carbon (million tCO₂e).]

V. DISCUSSION

The experimental results reveal several critical insights for AI governance. First, the inverse relationship between model capability and environmental efficiency necessitates explicit trade-off mechanisms in deployment decisions. While GPT-4 achieved superior ethical scores, its carbon footprint per inference was 2 $\times$ higher than Llama-3. This finding aligns with the UNU-INWEH report's conclusion that efficiency improvements alone cannot offset the rebound effect of increased usage [5].

Second, the geographic dimension of environmental impact emerges as a decisive factor. Shifting inference workloads from coal-dependent grids (carbon intensity: 800 gCO₂e/kWh) to renewable-rich regions (50 gCO₂e/kWh) reduces the carbon component of RAI by 93.75%, without hardware modifications. This validates the UN Secretary-General's call for renewable energy commitments and underscores the environmental justice implications of data center siting [13].

The geographic dependency of environmental impact, shown in Fig. 6, underscores a critical but underexplored dimension of AI sustainability: carbon intensity varies by a factor of 16 $\times$ across grids, from 50 gCO₂e/kWh in Iceland to 800 gCO₂e/kWh in coal-dependent regions. This disparity implies that infrastructure decisions---specifically data center siting---may exert greater influence on carbon footprint than algorithmic optimizations.




[Fig. 6. Carbon intensity ($\text{gCO}_2\text{e/kWh}$) and water factor (L/kWh) across regional electricity grids. Shifting inference from coal-dependent grids (China, India) to renewable-rich regions (Iceland) can reduce carbon emissions by over 90% without hardware modifications.]

Third, the framework's bias detection capabilities exposed systemic issues in military scenario outputs that existing fairness metrics overlooked. Within our simulated scenario evaluation, 34% of AI-generated policy recommendations exhibited subtle demographic biases undetectable by surface-level fairness metrics, underscoring the necessity of deep semantic analysis. Our Decision Risk Factor taxonomy provides a more granular approach, distinguishing between immediate physical harm (military) and cumulative distributive harm (policy).

Finally, the water footprint data challenges the carbon-centric narrative dominant in AI sustainability discourse. With AI systems potentially consuming more water annually than global bottled water production [18], and two-thirds of new data centers located in water-stressed regions [16], the framework's integrated water-energy-carbon analysis offers a more complete environmental picture than carbon-only assessments.

Fig. 7 presents the Pareto frontier between ethical and environmental performance, demonstrating that GPT-4 achieves superior ethical scores at the cost of higher environmental impact, while Llama-3 occupies the efficiency-dominant region---validating the framework's ability to expose these trade-offs quantitatively. The absence of any model in the ideal target zone ($E = 85, S = 85$) confirms that current architectures force an unavoidable compromise.



[Fig. 7. Pareto frontier between ethical performance (E) and environmental efficiency (S) for three LLMs. The gold star indicates the ideal target ($E = 85, S = 85$). No single model dominates both dimensions, necessitating explicit trade-off mechanisms in deployment decisions.]

VI. OPTIMIZATION FRAMEWORK

Based on experimental findings, we propose a three-tier optimization strategy:


1) Model Efficiency: Implement INT8 quantization and knowledge distillation to reduce parameter count while maintaining accuracy. Dynamic batching and speculative decoding can further reduce per-token energy consumption based on published efficiency studies [8].

2) Geographic Load Shifting: Deploy inference routing algorithms that prioritize low-carbon-intensity grids during peak demand. Temporal shifting (running non-urgent tasks during renewable surplus periods) can substantially reduce carbon footprint without latency penalties for asynchronous workloads [5][6].

3) Ethical Constraints in Prompt Engineering: Integrate constitutional AI principles directly into system prompts, reducing the need for extensive post-hoc bias correction. This approach, combined with retrieval-augmented generation (RAG) using curated, balanced datasets, improves both ethical scores and computational efficiency by reducing the need for extensive sampling.

VII. LIMITATIONS

As illustrated in Fig. 8, quantization emerges as the most favorable optimization strategy, achieving substantial energy reduction with negligible accuracy loss ($<1\%$), whereas knowledge distillation---while more aggressive---incurs greater performance penalties. Geographic load shifting offers unique advantages by decoupling computational efficiency from carbon intensity.



[Fig. 8. Optimization techniques: projected energy reduction versus estimated accuracy loss. Quantization (INT8) offers the highest energy savings with minimal accuracy degradation, while geographic load shifting provides carbon benefits without computational trade-offs.]

This work has several limitations that should be acknowledged. First, the experimental evaluation is simulation-based, using published hardware specifications and literature-derived empirical parameters rather than live GPU measurements. Consequently, absolute resource figures (e.g., 0.14 kWh per 1,000 inferences) should be interpreted as model-estimated values whose accuracy depends on the fidelity of the underlying assumptions; relative comparisons across models and scenarios remain more robust than absolute magnitudes.

Second, the RAI weighting coefficients ($w_1 = 0.4$, $w_2 = 0.35$, $w_3 = 0.25$) are currently grounded in literature alignment and regulatory priorities rather than empirically validated through structured user studies. While this approach ensures consistency with existing frameworks, the weights may require domain-specific calibration for deployment in healthcare, finance, or other specialized contexts.

Third, the framework was evaluated exclusively on text-generation tasks across three scenario categories (military, crisis, policy). Generalizability to multimodal systems (image, video, audio generation) and additional high-stakes domains (medical diagnosis, algorithmic trading) requires further validation, as environmental costs and ethical risks differ substantially across modalities.

Finally, the bias detection pipeline employs proxy metrics---demographic parity, equalized odds, and disparate impact---that may not capture intersectional or contextual forms of harm. Real-world deployment would benefit from integrating qualitative audits and stakeholder feedback mechanisms beyond automated scoring.

VIII. CONCLUSION AND FUTURE WORK

By integrating bias detection, explainability analysis, and estimated resource quantification into a single Responsible AI scoring model, the framework addresses a critical gap in current AI governance. Experimental validation across 150 simulated high-stakes scenarios demonstrates that the RAI score effectively discriminates between models along both ethical and sustainability axes, while optimization recommendations are projected to achieve substantial impact reductions without compromising accuracy.

The framework's alignment with emerging standards (ISO/IEC TR 20226:2025) and regulatory requirements (EU AI Act, UN Sustainable Development Goals) positions it as a practical tool for organizations seeking to deploy AI responsibly. The finding that inference---not

training---dominates lifecycle environmental costs (80-90% of total demand) redirects optimization efforts toward deployment efficiency, a shift with immediate practical implications.

Future work will extend the framework to: (1) incorporate land footprint metrics from critical mineral extraction for GPU manufacturing [20]; (2) develop real-time API integration for continuous monitoring; (3) expand scenario coverage to healthcare and financial domains; and (4) investigate the rebound effects of efficiency improvements on aggregate resource consumption. As Generative AI continues its exponential growth trajectory, frameworks like the one proposed here are essential to ensure that technological progress does not come at the cost of environmental degradation or ethical compromise.

IX. ACKNOWLEDGMENT

The author gratefully acknowledges the mentorship of Mohammad Sameer (Generative AI Developer, Phinerd) and the support of JNTU-GV College of Engineering, Vizianagaram. This research was conducted as part of an undergraduate research initiative in Artificial Intelligence and Sustainable Computing.

X. DATA AND CODE AVAILABILITY

The framework code, simulation modules, scenario datasets, and visualization scripts are available as open-source at <https://github.com/GSMS-B/ethical-sustainable-genai-framework>. The paper uses publicly available literature-derived parameters whose sources are cited in the References. All simulation scenarios, ethical evaluation prompts, and environmental estimation parameters are documented in the repository to ensure full reproducibility.

REFERENCES

- [1] OpenAI, 'GPT-4 Technical Report,' arXiv:2303.08774, 2023.
- [2] S. Russell, 'Human Compatible: Artificial Intelligence and the Problem of Control,' Viking, 2019.
- [3] E. Strubell, A. Ganesh, and A. McCallum, 'Energy and Policy Considerations for Deep Learning in NLP,' Proc. of the 57th ACL, pp. 3645-3650, 2019.
- [4] A. de Vries, 'The Growing Energy Footprint of Artificial Intelligence,' Joule, vol. 8, no. 6, pp. 1-4, 2025.
- [5] M. Aczel et al., 'Environmental Cost of Artificial Intelligence: Carbon, Water, and Land Footprints,' UNU-INWEH Report, doi: 10.53328/INR26RMA002, 2026.
- [6] International Energy Agency (IEA), 'Electricity 2025: Analysis and Forecast to 2027,' Paris, 2025.
- [7] E. M. Bender et al., 'On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?' Proc. of FAccT, pp. 610-623, 2021.
- [8] D. Patterson et al., 'Carbon Emissions and Large Neural Network Training,' arXiv:2104.10350, 2021.
- [9] N. Mehrabi et al., 'A Survey on Bias and Fairness in Machine Learning,' ACM Computing Surveys, vol. 54, no. 6, 2021.
- [10] A. B. Arrieta et al., 'Explainable Artificial Intelligence (XAI): Concepts, Taxonomies, Opportunities and Challenges,' Information Fusion, vol. 58, pp. 82-115, 2020.
- [11] European Parliament, 'Artificial Intelligence Act,' Regulation (EU) 2024/1689, 2024.
- [12] ISO/IEC, 'TR 20226:2025 Information Technology --- Artificial Intelligence --- Environmental Sustainability Aspects,' 2025.
- [13] A. Guterres, 'Remarks at London Climate Action Week,' United Nations, June 2025.
- [14] P. Li et al., 'Making AI Less 'Thirsty': Uncovering and Addressing the Secret Water Footprint of AI Models,' arXiv:2304.03271, 2023.
- [15] S. Luccioni et al., 'Estimating the Carbon Footprint of BLOOM, a 176B Parameter Language Model,' Journal of Machine Learning Research, vol. 24, pp. 1-15, 2023.
- [16] ANSI, 'Making AI Data Centers More Sustainable: Using Standards,' ANSI Blog, May 2026.
- [17] B. Li et al., 'Carbon in Motion: Characterizing Open-Sora on the Sustainability of Generative AI for Video Generation,' Proc. of the ACM HotCarbon Workshop, 2024.
- [18] A. de Vries, 'The Carbon and Water Footprints of Data Centers,' Patterns, vol. 6, no. 3, 2025.
- [19] M. S. Islam et al., 'Green AI: A Systematic Literature Review,' IEEE Access, vol. 12, pp. 34567-34589, 2024.
- [20] Y. Bengio et al., 'Managing AI Risks in an Era of Rapid Progress,' Science, vol. 384, no. 6698, pp. 842-844, 2024.
- [21] S. Bommasani et al., 'On the Opportunities and Risks of Foundation Models,' arXiv:2108.07258, 2021.

AUTHOR BIOGRAPHY

Bugatha Ganasyam Mani Sankar is a final year B.Tech student in Computer Science and Engineering at JNTU-GV College of Engineering, Vizianagaram, India. His research interests include Generative AI, AI Ethics, Green Computing, and Sustainable Machine Learning. He is currently working under the mentorship of Mohammad Sameer at Phinerd on responsible AI frameworks.